

# uC/OS-II Part 7: Mutual Exclusion Semaphores

Real-Time Compting  
Prof. Li-Pin Chang  
ESSLab@NCTU

# Mutual Exclusion Semaphores

- A mutex is used to synchronize tasks with priority inversion management
- If a LPT blocks a HPT, the priority of the LPT is boosted
  - Priority inheritance
  - With priority inheritance, multiple tasks might have the same priority

# Mutual Exclusion Semaphores

- uC/OS-II uses an alternative approach to work around the priority duplication issue
  - Reserve a priority for a mutex
  - The reserved priority (for the mutex) must be higher than all tasks that use the mutex

# Mutual Exclusion Semaphores

- Be careful about the priority for mutex!!
  - It should be immediately higher than the highest priority of all tasks use the mutex
  - To avoid affecting unrelated tasks

# Example1

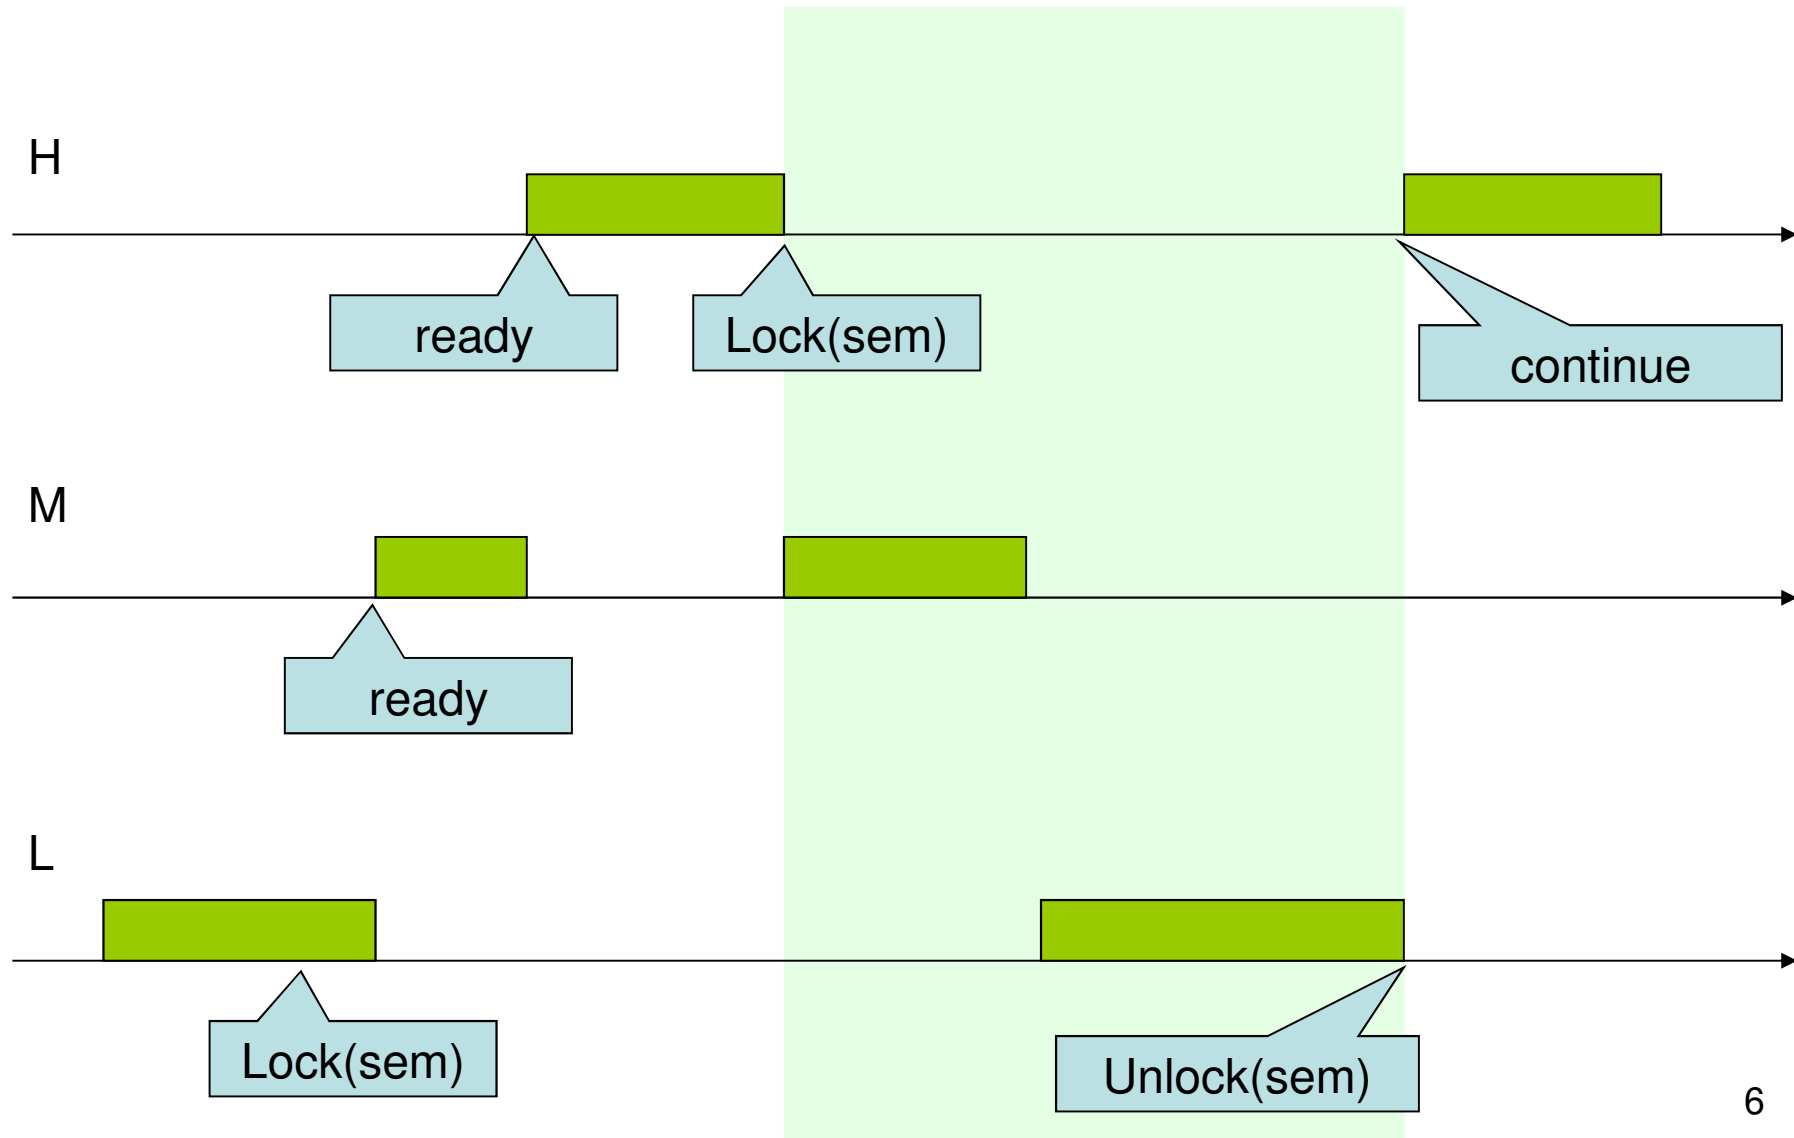
```
OSMutexCreate(VH, &err);

void taskPrioH {
    while(1) {
        /*...*/
        OSMutexPend(mutex, 0, &err);
        /*...*/
        OSMutexPost(Mutex);
    }
}
```

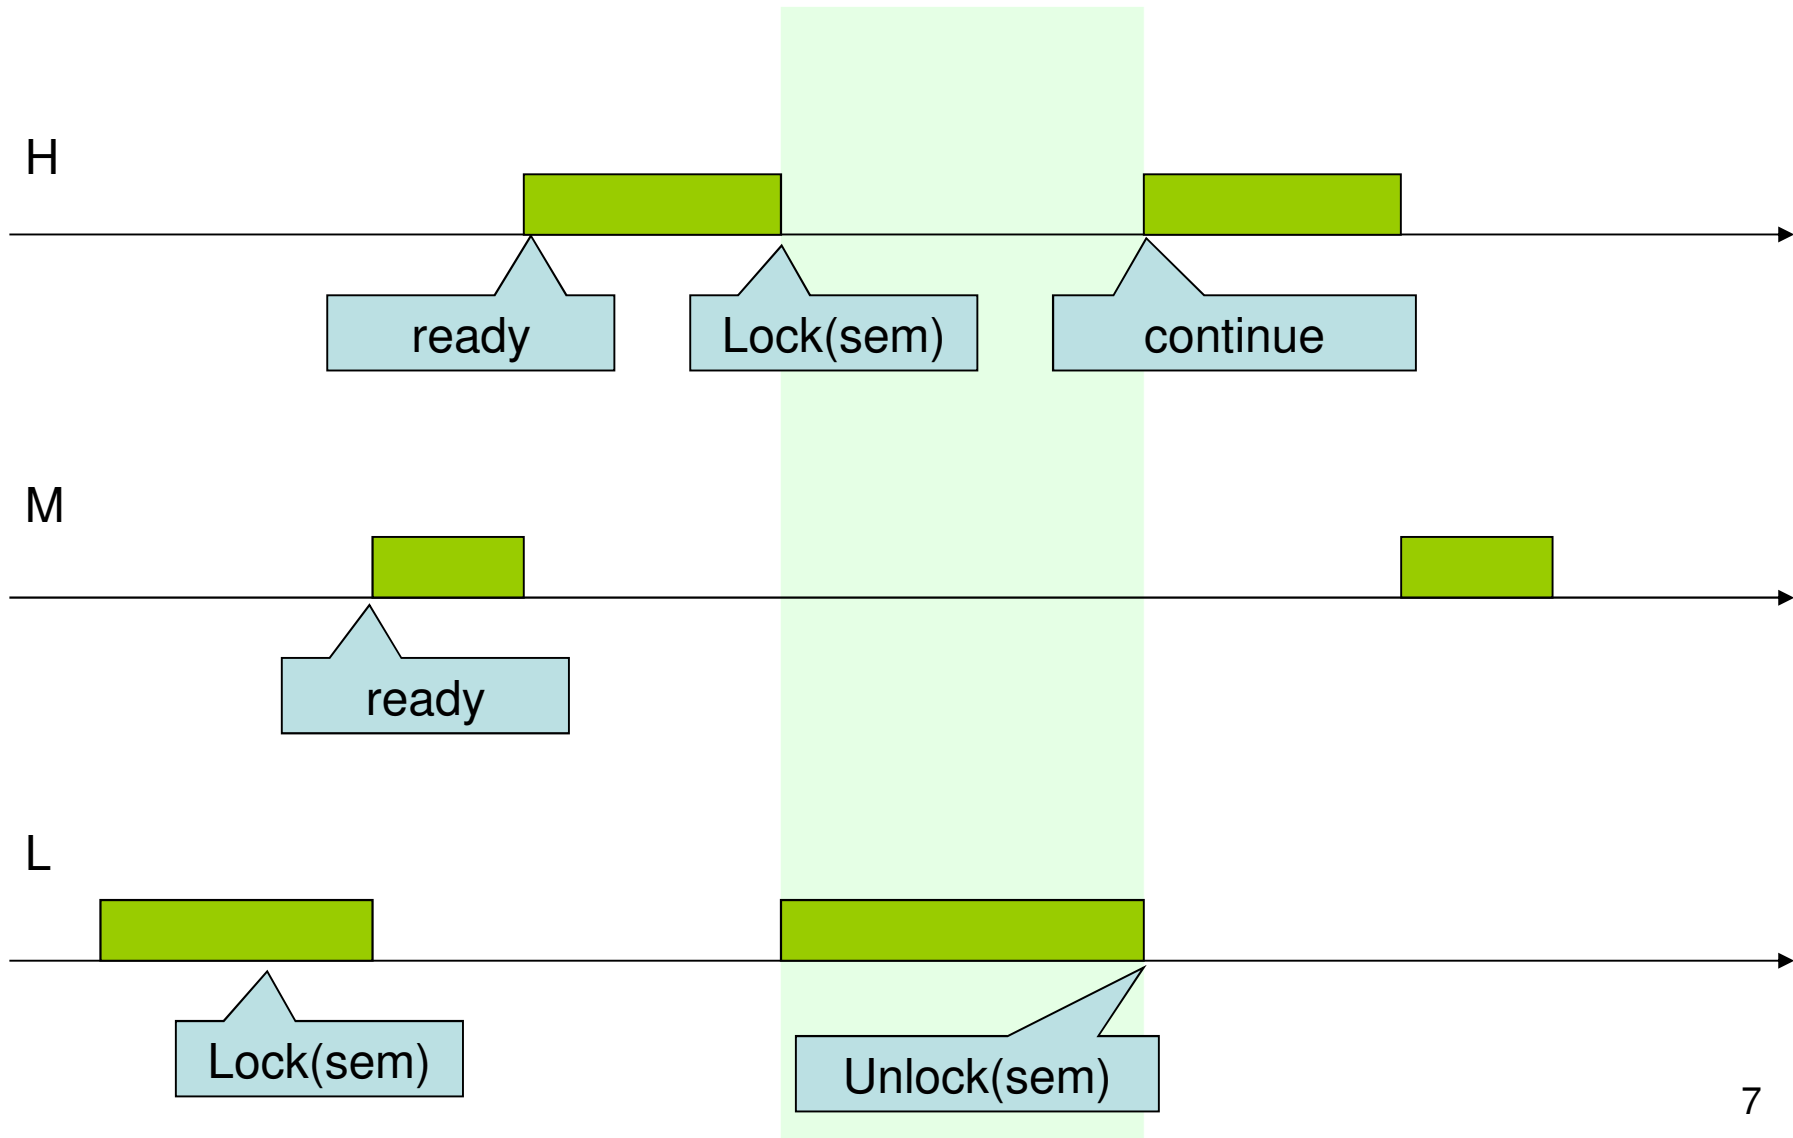
```
void taskPrioM {
    while(1) {
        /*...*/
    }
}
```

```
void taskPrioL {
    while(1) {
        /*...*/
        OSMutexPend(mutex, 0, &err);
        /*...*/
        OSMutexPost(Mutex);
    }
}
```

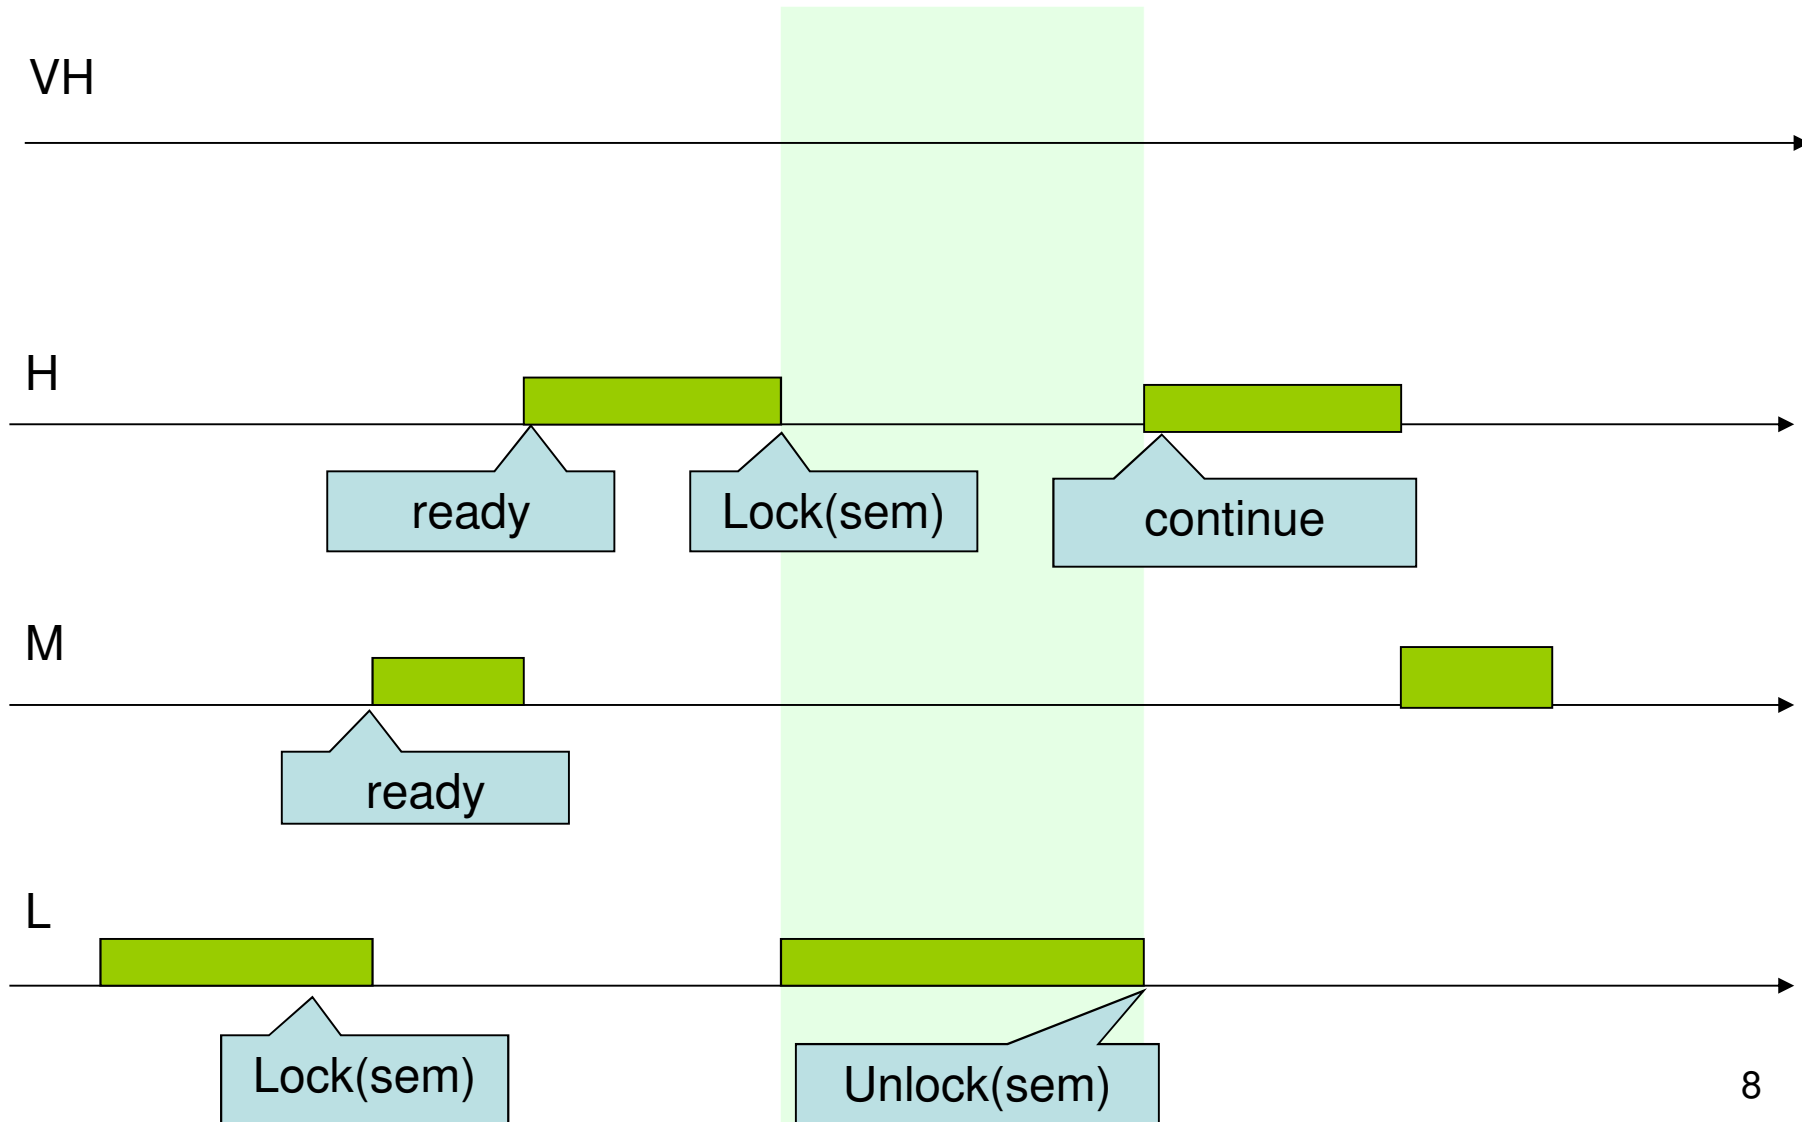
# Example - without PIP



# Example - with PIP



# Example - $\mu$ C/OS-II





# Mutual Exclusion Semaphores

- uC/OS-II
  - It differs from NPCS
    - The response is not that bad (not non-preemptible)
  - It differs from ceiling priority protocol (CPP)
    - CPP raises priority when resource lock is **acquired**
    - uC/OS-II raises priority when LPT **blocks** HPT
  - It differs from priority-inheritance protocol (PIP)
    - uc/OS-II has no incremental priority inheritance

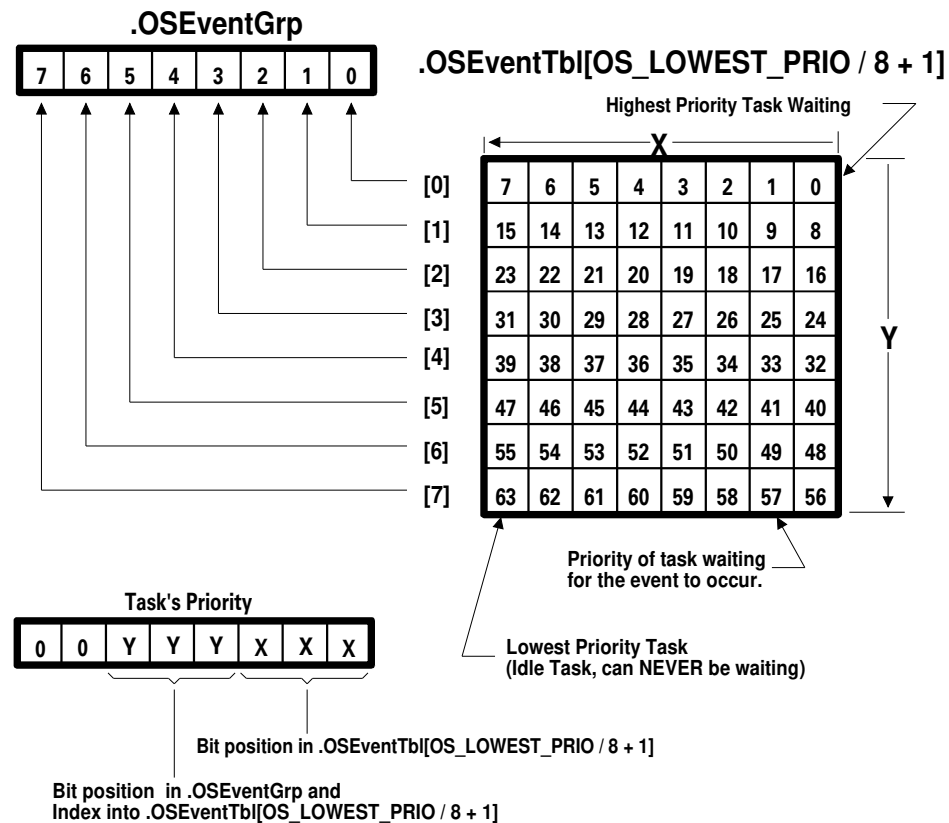
# Functions

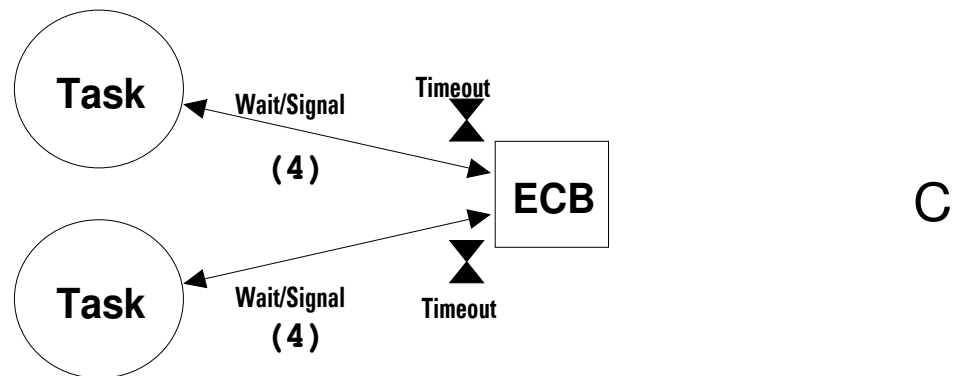
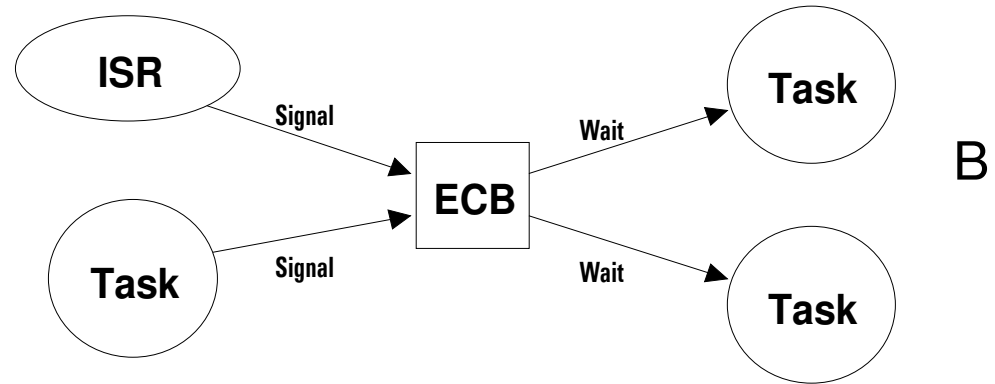


- OSMutexCreate()
- OSMutexPend()
- OSMutexPost()
- OSMutexDel()
- OSMutexAccept()
- OSMutexQuery()

```

typedef struct {
    INT8U  OSEventType;           /* Event type */
    INT8U  OSEventGrp;           /* Group for wait list */
    INT16U OSEventCnt;           /* Count (when event is a semaphore) */
    void  *OSEventPtr;           /* Ptr to message or queue structure */
    INT8U  OSEventTbl[OS_EVENT_TBL_SIZE]; /* Wait list for event to occur */
} OS_EVENT;

```





No pending in ISR

# OSMutexCreate()

- OSMutexCreate() stores the PIP priority in OSEventCnt
  - OSEventCnt stores a value if the event is a semaphore

Type	OS_EVENT_TYPE_MUTEX	
Grp	0x00	
Cnt	PIP prio	0xFF
Ptr	NULL	
Tbl	0x00 ... 0x00	

available

OS_EVENT_TYPE_MUTEX	
???	
PIP prio	Owner task priority
.....	
??? ... ???	

Owned by a task

TCB of  
mutex  
owner

```

OS_EVENT *OSMutexCreate (INT8U prio, INT8U *err)
{
    #if OS_CRITICAL_METHOD == 3                                /* Allocate storage for CPU status register */
        OS_CPU_SR cpu_sr;
    #endif
        OS_EVENT *pevent;

    if (OSIntNesting > 0) {                                    /* See if called from ISR ... */
        *err = OS_ERR_CREATE_ISR;                             /* ... can't CREATE mutex from an ISR */
        return ((OS_EVENT *)0);
    }
    #if OS_ARG_CHK_EN > 0
        if (prio >= OS_LOWEST_PRIO) {                          /* Validate PIP */
            *err = OS_PRIO_INVALID;
            return ((OS_EVENT *)0);
        }
    #endif
    OS_ENTER_CRITICAL();
    if (OSTCBPrioTbl[prio] != (OS_TCB *)0) {                  /* Mutex priority must not already exist */
        OS_EXIT_CRITICAL();                                    /* Task already exist at priority ... */
        *err = OS_PRIO_EXIST;                                  /* ... inheritance priority */
        return ((OS_EVENT *)0);
    }
    OSTCBPrioTbl[prio] = (OS_TCB *)1;                          /* Reserve the table entry */
    pevent = OSEventFreeList;                                  /* Get next free event control block */
    if (pevent == (OS_EVENT *)0) {                             /* See if an ECB was available */
        OSTCBPrioTbl[prio] = (OS_TCB *)0;                     /* No, Release the table entry */
        OS_EXIT_CRITICAL();
        *err = OS_ERR_PEVENT_NULL;                             /* No more event control blocks */
        return (pevent);
    }
    OSEventFreeList = (OS_EVENT *)OSEventFreeList->OSEventPtr; /* Adjust the free list */
    OS_EXIT_CRITICAL();
    pevent->OSEventType = OS_EVENT_TYPE_MUTEX;
    pevent->OSEventCnt = (prio << 8) | OS_MUTEX_AVAILABLE; /* Resource is available */
    pevent->OSEventPtr = (void *)0;                             /* No task owning the mutex */
    OS_EventWaitListInit(pevent);
    *err = OS_NO_ERR;
    return (pevent);
}

```

Reserve a priority by placing a non-null value

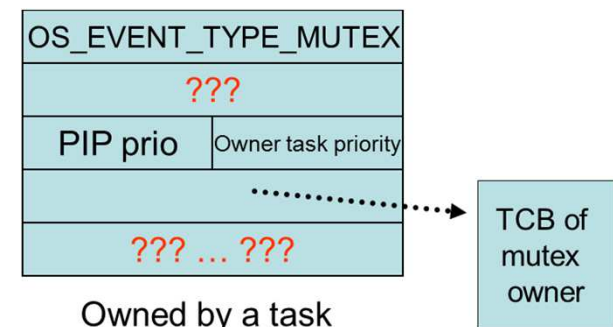
# OSMutexPend()

```

void OSMutexPend (OS_EVENT *pevent, INT16U timeout, INT8U *err) {
    INT8U    pip;                                /* Priority Inheritance Priority (PIP) */
    INT8U    mprio;                              /* Mutex owner priority */
    BOOLEAN   rdy;                               /* Flag indicating task was ready */
    OS_TCB    *ptcb;
    if (OSIntNesting > 0) {                      /* See if called from ISR ... */
        *err = OS_ERR_PEND_ISR;                 /* ... can't PEND from an ISR */
        return;
    }
    OS_ENTER_CRITICAL();                          /* Is Mutex available? */
    if ((INT8U) (pevent->OSEventCnt & OS_MUTEX_KEEP_LOWER_8) == OS_MUTEX_AVAILABLE) {
        pevent->OSEventCnt &= OS_MUTEX_KEEP_UPPER_8; /* Yes, Acquire the resource */
        pevent->OSEventCnt |= OSTCBCur->OSTCBPrio;    /* Save priority of owning task */
        pevent->OSEventPtr = (void *) OSTCBCur;      /* Point to owning task's OS_TCB */
        OS_EXIT_CRITICAL();
        *err = OS_NO_ERR;
        return;
    }
    pip = (INT8U) (pevent->OSEventCnt >> 8);        /* No, Get PIP from mutex */
    mprio = (INT8U) (pevent->OSEventCnt & OS_MUTEX_KEEP_LOWER_8); /* Get priority of mutex owner */
    ptcb = (OS_TCB *) (pevent->OSEventPtr);          /* Point to TCB of mutex owner */
}

```

If the mutex is free...



```

if (ptcb->OSTCBPrio != pip && mprio > OSTCBCur->OSTCBPrio) { /* Need to promote prio of owner?*/
    if ((OSRdyTbl[ptcb->OSTCBBY] & ptcb->OSTCBBitX) != 0x00) { /* See if mutex owner is ready */
                                                                    /* Yes, Remove owner from Rdy ...*/
                                                                    /* ... list at current prio */
        if ((OSRdyTbl[ptcb->OSTCBBY] & ~ptcb->OSTCBBitX) == 0x00) {
            OSRdyGrp &= ~ptcb->OSTCBBitY;
        }
        rdy = TRUE;
    } else {
        rdy = FALSE;
    }
    ptcb->OSTCBPrio = pip; /* Change owner task prio to PIP */
    ptcb->OSTCBBY = ptcb->OSTCBPrio >> 3;
    ptcb->OSTCBBitY = OSMapTbl[ptcb->OSTCBBY];
    ptcb->OSTCBBitX = ptcb->OSTCBPrio & 0x07;
    ptcb->OSTCBBitX = OSMapTbl[ptcb->OSTCBBitY];
    if (rdy == TRUE) { /* If task was ready at owner's priority ...*/
        OSRdyGrp |= ptcb->OSTCBBitY; /* ... make it ready at new priority. */
        OSRdyTbl[ptcb->OSTCBBY] |= ptcb->OSTCBBitX;
    }
    ★ OSTCBPrioTbl[pip] = (OS_TCB *)ptcb;
}
OSTCBCur->OSTCBStat |= OS_STAT_MUTEX; /* Mutex not available, pend current task */
OSTCBCur->OSTCBDly = timeout; /* Store timeout in current task's TCB */
OS_EventTaskWait(pevent); /* Suspend task until event or timeout occurs */
OS_EXIT_CRITICAL();
★ OS_Sched(); /* Find next highest priority task ready */
OS_ENTER_CRITICAL();
if (OSTCBCur->OSTCBStat & OS_STAT_MUTEX) { /* Must have timed out if still waiting for event*/
    OS_EventTO(pevent);
    OS_EXIT_CRITICAL();
    *err = OS_TIMEOUT; /* Indicate that we didn't get mutex within TO */
    return;
}
OSTCBCur->OSTCBEventPtr = (OS_EVENT *)0;
OS_EXIT_CRITICAL();
*err = OS_NO_ERR;
}

```

- If the owner's priority has not been raised, do it.
- If the owner's is current ready, change its bit from the ready-list bitmap

Set the ECB bitmap,  
clear itself from RdyMap

Locked the mutex without timed-out

pip = mutex's PIP priority  
mprio = mutex owner's priority (original)  
ptcb->OSTCBPrio = mutex owner's priority (current)



# OSMutexPost()

```
INT8U OSMutexPost (OS_EVENT *pevent) {  
    INT8U    pip;                                /* Priority inheritance priority */  
    INT8U    prio;                                /*  
    if (OSIntNesting > 0) {                       /* See if called from ISR ... */  
        return (OS_ERR_POST_ISR);                 /* ... can't POST mutex from an ISR */  
    }  
    OS_ENTER_CRITICAL();  
    pip = (INT8U) (pevent->OSEventCnt >> 8);      /* Get priority inheritance priority of mutex */  
    prio = (INT8U) (pevent->OSEventCnt & OS_MUTEX_KEEP_LOWER_8); /* Get owner's original priority */  
    if (OSTCBCur->OSTCBPrio != pip &&             /* See if posting task owns the MUTEX */  
        OSTCBCur->OSTCBPrio != prio) {  
        OS_EXIT_CRITICAL();  
        return (OS_ERR_NOT_MUTEX_OWNER);  
    }  
}
```

The task posting the mutex must be the task that owns (has acquired) the mutex!!

# OSMutexPost()

```

if (OSTCBCur->OSTCBPrio == pip) {
    /* Did we have to raise current task's priority? */
    /* Yes, Return to original priority */
    /* Remove owner from ready list at 'pip' */
    if ((OSRdyTbl[OSTCBCur->OSTCBY] &= ~OSTCBCur->OSTCBBitX) == 0) {
        OSRdyGrp &= ~OSTCBCur->OSTCBBitY;
    }
    OSTCBCur->OSTCBPrio = prio;
    OSTCBCur->OSTCBY = prio >> 3;
    OSTCBCur->OSTCBBitY = OSMaTbl[OSTCBCur->OSTCBY];
    OSTCBCur->OSTCBX = prio & 0x07;
    OSTCBCur->OSTCBBitX = OSMaTbl[OSTCBCur->OSTCBX];
    OSRdyGrp |= OSTCBCur->OSTCBBitY;
    OSRdyTbl[OSTCBCur->OSTCBY] |= OSTCBCur->OSTCBBitX;
    ★ OSTCBPrioTbl[prio] = (OS_TCB *)OSTCBCur;
}
OSTCBPrioTbl[pip] = (OS_TCB *)1;
if (pevent->OSEventGrp != 0x00) {
    ★ prio = OS_EventTaskRdy(pevent, (void *)0, OS_STAT_MUTEX);
    pevent->OSEventCnt &= OS_MUTEX_KEEP_UPPER_8;
    pevent->OSEventCnt |= prio;
    pevent->OSEventPtr = OSTCBPrioTbl[prio];
    OS_EXIT_CRITICAL();
    ★ OS_Sched();
    return (OS_NO_ERR);
}
pevent->OSEventCnt |= OS_MUTEX_AVAILABLE;
pevent->OSEventPtr = (void *)0;
OS_EXIT_CRITICAL();
return (OS_NO_ERR);
}

```

Move the ready bit of the current task back to its original position

Release the highest waiting task and transfer the ownership of the mutex to the task

# OSMutexDel()

```
OS_EVENT *OSMutexDel (OS_EVENT *pevent, INT8U opt, INT8U *err) {
    BOOLEAN    tasks_waiting;
    INT8U      pip;
    if (OSIntNesting > 0) {
        /* See if called from ISR ... */
        /* ... can't DELETE from an ISR */
        *err = OS_ERR_DEL_ISR;
        return (pevent);
    }
    OS_ENTER_CRITICAL();
    if (pevent->OSEventGrp != 0x00) {
        /* See if any tasks waiting on mutex */
        /* Yes */
        tasks_waiting = TRUE;
    } else {
        /* No */
        tasks_waiting = FALSE;
    }

    switch (opt) {
        case OS_DEL_NO_PEND:
            /* Delete mutex only if no task waiting */
            if (tasks_waiting == FALSE) {
                pip = (INT8U) (pevent->OSEventCnt >> 8);
                OSTCBPrioTbl[pip] = (OS_TCB *)0; /* Free up the PIP */
                pevent->OSEventType = OS_EVENT_TYPE_UNUSED;
                pevent->OSEventPtr = OSEventFreeList; /* Return Event Control Block to free list */
                OSEventFreeList = pevent;
                OS_EXIT_CRITICAL();
                *err = OS_NO_ERR;
                return ((OS_EVENT *)0); /* Mutex has been deleted */
            } else {
                OS_EXIT_CRITICAL();
                *err = OS_ERR_TASK_WAITING;
                return (pevent);
            }
    }
}
```

# OSMutexDel()

```
case OS_DEL_ALWAYS:                                     /* Always delete the mutex */
    while (pevent->OSEventGrp != 0x00) {                 /* Ready ALL tasks waiting for mutex */
        OS_EventTaskRdy(pevent, (void *)0, OS_STAT_MUTEX);
    }
    pip = (INT8U) (pevent->OSEventCnt >> 8);
    OSTCBPrioTbl[pip] = (OS_TCB *)0;                     /* Free up the PIP */
    pevent->OSEventType = OS_EVENT_TYPE_UNUSED;
    pevent->OSEventPtr = OSEventFreeList;                 /* Return Event Control Block to free list */
    OSEventFreeList = pevent;                             /* Get next free event control block */
    OS_EXIT_CRITICAL();
    if (tasks_waiting == TRUE) {                          /* Reschedule only if task(s) were waiting */
        OS_Sched();                                       /* Find highest priority task ready to run */
    }
    *err = OS_NO_ERR;
    return ((OS_EVENT *)0);                               /* Mutex has been deleted */

default:
    OS_EXIT_CRITICAL();
    *err = OS_ERR_INVALID_OPT;
    return (pevent);
}
}
```

# OSMutexAccept()

```
INT8U OSMutexAccept (OS_EVENT *pevent, INT8U *err) {
    if (OSIntNesting > 0) {                                /* Make sure it's not called from an ISR */
        *err = OS_ERR_PEND_ISR;
        return (0);
    }
    OS_ENTER_CRITICAL();                                    /* Get value (0 or 1) of Mutex*/
    if ((pevent->OSEventCnt & OS_MUTEX_KEEP_LOWER_8) == OS_MUTEX_AVAILABLE) {
        pevent->OSEventCnt &= OS_MUTEX_KEEP_UPPER_8;      /* Mask off LSByte (Acquire Mutex) */
        pevent->OSEventCnt |= OSTCBCur->OSTCBPrio;         /* Save current task priority in LSByte */
        pevent->OSEventPtr = (void *) OSTCBCur;           /* Link TCB of task owning Mutex */
        OS_EXIT_CRITICAL();
        *err = OS_NO_ERR;
        return (1);
    }
    OS_EXIT_CRITICAL();
    *err = OS_NO_ERR;
    return (0);
}
```

# Summary

- In realistic systems, compromise exists between simplicity and performance
  - [PCP]→[PIP]→[CPP]→[NPCS]
- With uc/OS-II's mutex, how is
  - Priority inversion management?
  - deadlock avoidance?